

# problems4

November 26, 2018

## 1 Problem sheet 4

### 1.1 1

In this problem, we generate Nietzsche text based on embedding words instead of characters. We adapt the keras-code from [https://keras.rstudio.com/articles/examples/pretrained\\_word\\_embeddings.html](https://keras.rstudio.com/articles/examples/pretrained_word_embeddings.html).

Before getting started, think about possible advantages and disadvantages of using word embeddings

Next, download the [Glove word vectors](#) and create a look-up table.

```
In [ ]: emb_idx <- new.env(parent = emptyenv())
        lines <- readLines("./data/glove.6B.100d.txt")

        for (line in lines) {
          values <- strsplit(line, ' ', fixed = TRUE)[[1]]
          w <- values[[1]]
          coefs <- as.numeric(values[-1])
          emb_idx[[w]] <- coefs
        }
```

Proceed as in the lecture to load the Nietzsche text into a variable text. Then, tokenize the text.

```
In [ ]: tokenizer <- text_tokenizer()
        tokenizer %>% fit_text_tokenizer(text)
        t2s<- texts_to_sequences(tokenizer, text)[[1]]
```

Now, that the words are tokenized, we can set up the embedding matrix.

```
In [ ]: emb_dim <- 100
        w_idx <- tokenizer$word_index
        nwords <- length(w_idx) + 1

        emb_mat <- matrix(0L, nrow = nwords, ncol = emb_dim)
        for (w in names(w_idx)){
          idx <- w_idx[[w]]
          emb_vec <- emb_idx[[w]]
```

```

# words not found in embedding index will be all-zeros.
if (!is.null(emb_vec))
  emb_mat[index,] <- emb_vec
}

```

The embedding matrix now becomes part of an embedding layer.

```

In [ ]: emb_layer <- layer_embedding(
  input_dim = nwords,
  output_dim = emb_dim,
  weights = list(emb_mat),
  input_length = maxlen,
  trainable = F
)

```

Build a keras model based on the embedding layer.

```

In [ ]: input <- layer_input(shape = list(maxlen),
                             dtype = 'int32')
output <- ??

model <- keras_model(input, output)
model %>% compile(
  loss = "sparse_categorical_crossentropy",
  optimizer = optimizer_rmsprop(.01)
)

```

Finally, we build the training data.

As in the lecture, chop the sequence into smaller bits and collect them in a variable `sentences`. Also store the next word into `next_words`.

```

In [ ]: sentences <- ??
next_words <- ??

```

Fit the model!

## 1.2 2

Modify the weather example from the lecture to incorporate validation data.

## 1.3 3

Develop the [seq2seq architecture](#) for the weather example.

The idea is to use two neural networks, an **encoder** and a **decoder**. The encoder analyzes the history and summarizes it into a single vector. The decoder takes this vector and generates from it a target sequence.

<https://arxiv.org/abs/1409.3215v3>

In a wonderful [blog post](#), François Chollet explains how to put the seq2seq-architecture into 'keras'. The corresponding pattern for R is described in <https://github.com/rstudio/keras/issues/312>. First, the encoder is constructed as

```
In [ ]: encoder_inputs <- layer_input(shape = list(NULL, dim_data))
        encoder_results <- encoder_inputs %>%
          layer_lstm(units = cell_size, return_state = T)
        encoder_states <- encoder_results[2:3]
```

Second, the decoder reads as

```
In [ ]: decoder_inputs <- layer_input(shape = list(NULL, dim_pred))

        decoder_lstm <- layer_lstm(units = cell_size,
                                     return_sequences = T,
                                     return_state = T,
                                     stateful = F)
        decoder_results <- decoder_lstm(decoder_inputs, initial_state=encoder_states)
        decoder_dense <- layer_dense(dim_pred, activation=?)
        decoder_outputs <- decoder_dense(decoder_results[[1]])
```

Now, adapt this idea to the weather data.